

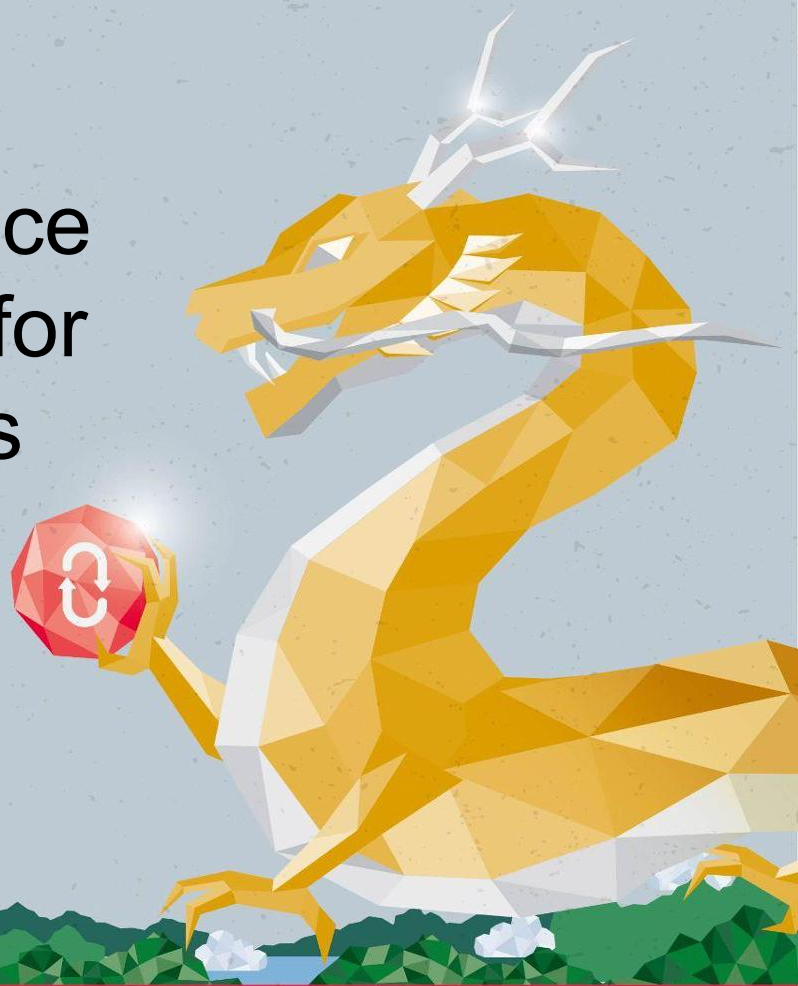
CURVE: a High-performance Software-Defined storage for Cloud Native applications



Curve

Xiang Dong

Architect, Netease Hangzhou Research Institute



Agenda

✓ Why develop storage

Design objectives

Achievements in CURVE

Key designs used by CURVE

CURVE Roadmap

CURVE Community

You can take away



Why develop

Storage requirements

- Hardware requirements
 - more CPU cores, faster Network, nvme storage
- Problems for stateful apps
 - storage capacity expansion
 - capacity imbalance
 - apps bundled with data locations
- Requirements for elastic block storage
- Requirements for file system



open-source storage

- Requirements
 - Cloud Native
 - Easy operation and maintenance
 - High performance
- CEPH
 - Complex and Large amount of codes
 - system maintenance are difficult in corner cases
 - Performance requirements cannot be met



Agenda

Why develop storage

✓ Design objectives

Achievements in CURVE

Key designs used by CURVE

CURVE Roadmap

CURVE Community

You can take away



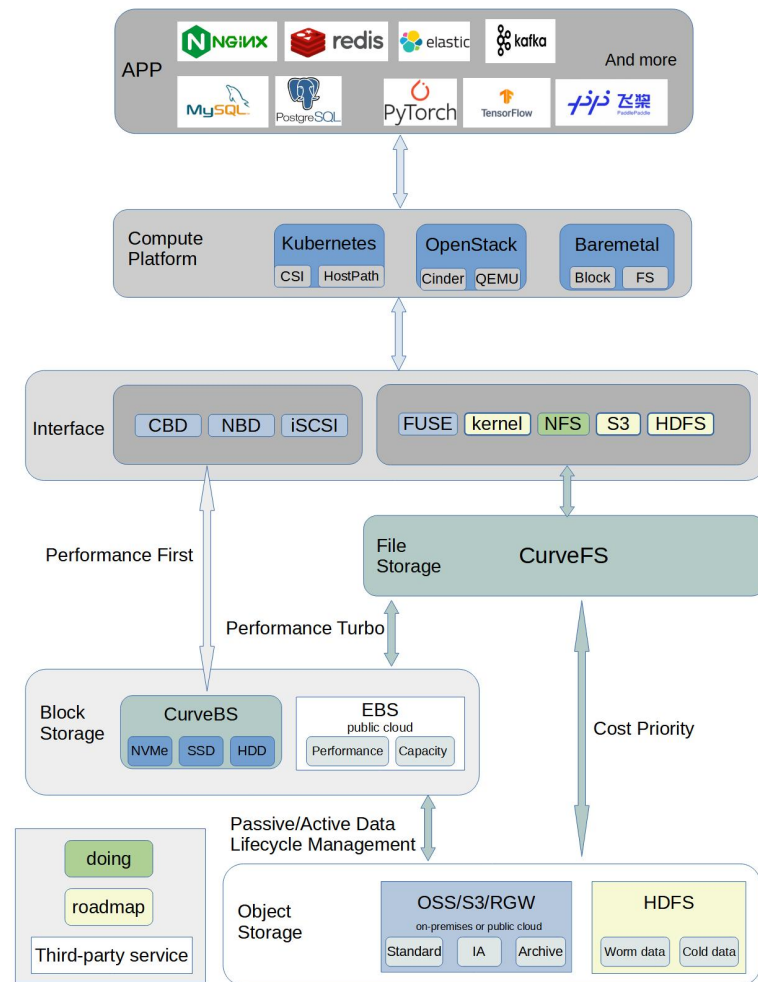
Design objectives

App scenario & requires

- Database scenario
- Elastic block Storage for KVM / Kubernetes / iSCSI

Design objectives

- Thin provisioning storage pools
- High performance
- Easy maintenance
- Cloud Native



Agenda

Why develop storage

Design objectives

✓ Achievements in CURVE

Key designs used by CURVE

CURVE Roadmap

CURVE Community

You can take away



Performance in database scenario

Solution architecture

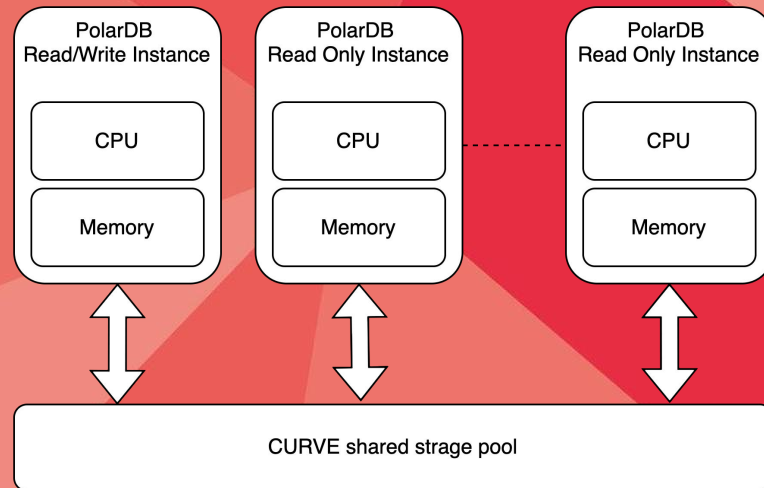
- CURVE provides underlying distributed shared storage services
- PolarDB for PostgreSQL provides database services

Solution features

- separation between storage and computation
- The call method is the same as that of the standalone database

Performance test vs CEPH

- BenchmarkSQL transactions per minute increased by 39%
- In PGbench tests, latency decreased by 21% and TPS improved by 26%



Performance compare with CEPH

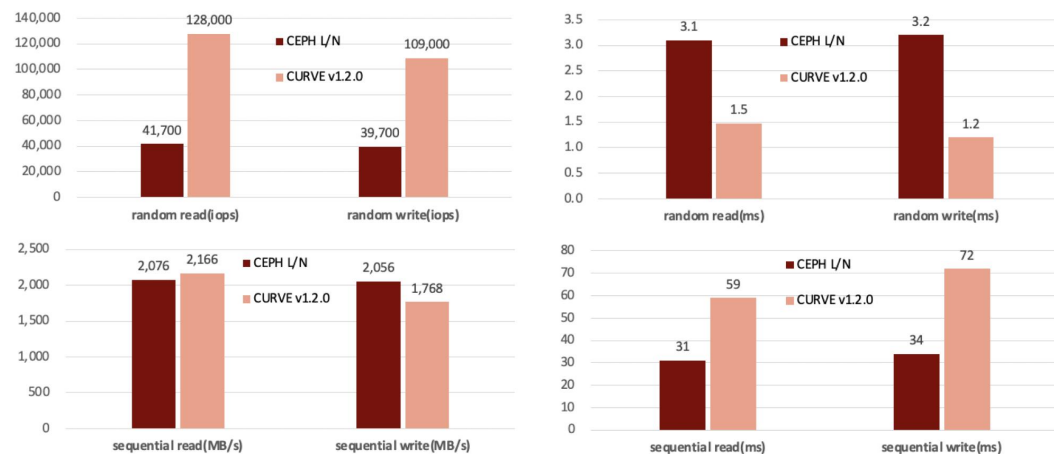
The test environment

- A cluster consists of six nodes.
Each node consists of 20 x SSDS,
256GB memory, and 2 x e5-2660
cpus

Performance using RDMA

- Compared with TCP, I/O throughput
increases by 20%

vs of Single Vol Environment: 3 replicas on a 6 nodes cluster, each node has 20xSATA SSD,
2xE5-2660 v4 and 256GB memory



Availability testing in corner cases

The impact of different faults on I/O jitter delay

- Disk failure
- Server failure
- Server fake death
- Slow response

Faults case	curve i/o 抖动
Disk failure	4s
Server failure	4s
Server fake death	4s
Slow response	1s frequently

Data availability analysis

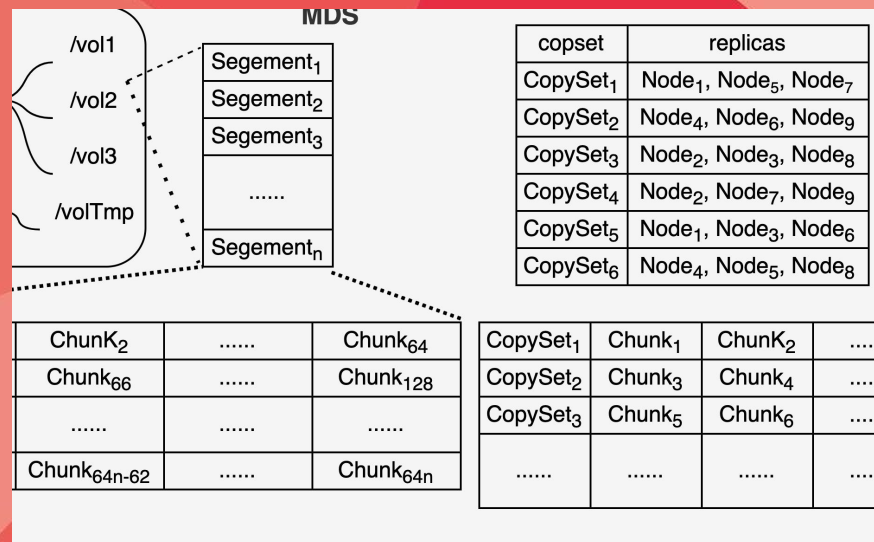
Copyset allocation algorithm

cluster with 1200 disks (each have MTBF 1.2 million hours)

3 replicas

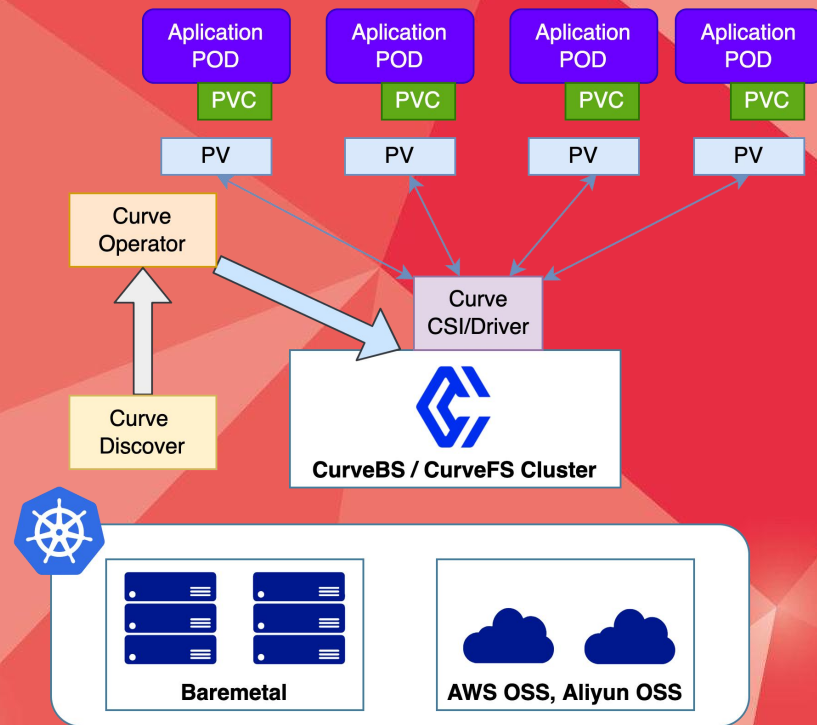
Restore data on a disk within 5 minutes

Data availability of 6 nines can be achieved



Cloud native Support

Currently we offer CSI Driver for block storage to provide PV/PVC resources on Kubernetes



Agenda

Why develop storage

Design objectives

Achievements in CURVE

✓ Key designs used by CURVE

CURVE Roadmap

CURVE Community

You can take away



Key designs used by CURVE

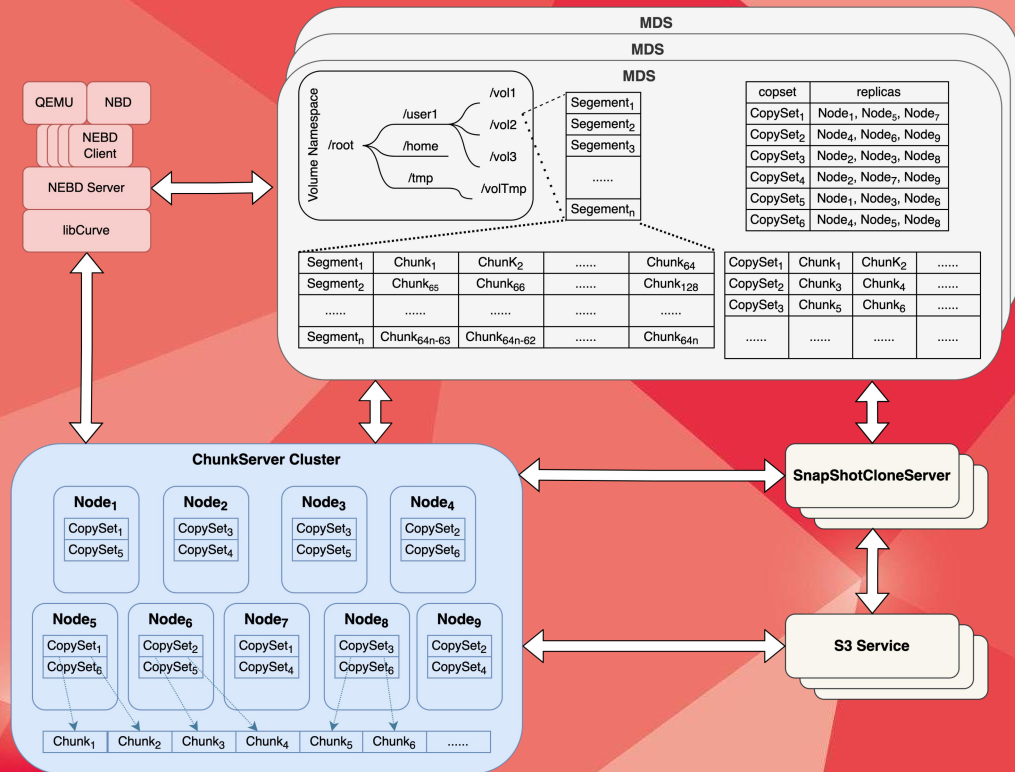
High availability and reliability

- Cluster topology
- CopySet pre-allocation algorithm
- Raft Consistency protocol

High performance

- pre-created file pool
- data strip like RAID
- Zero data copy
- RDMA

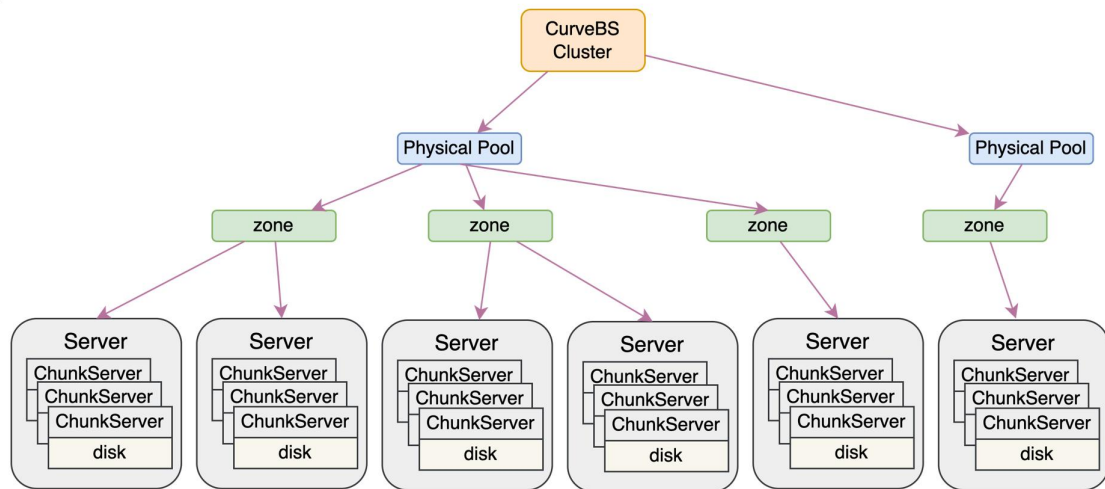
Cloud Native



Zone is the basic unit of fault isolation

Server Indicates a physical server

Chunkserver is a service instance on a physical server



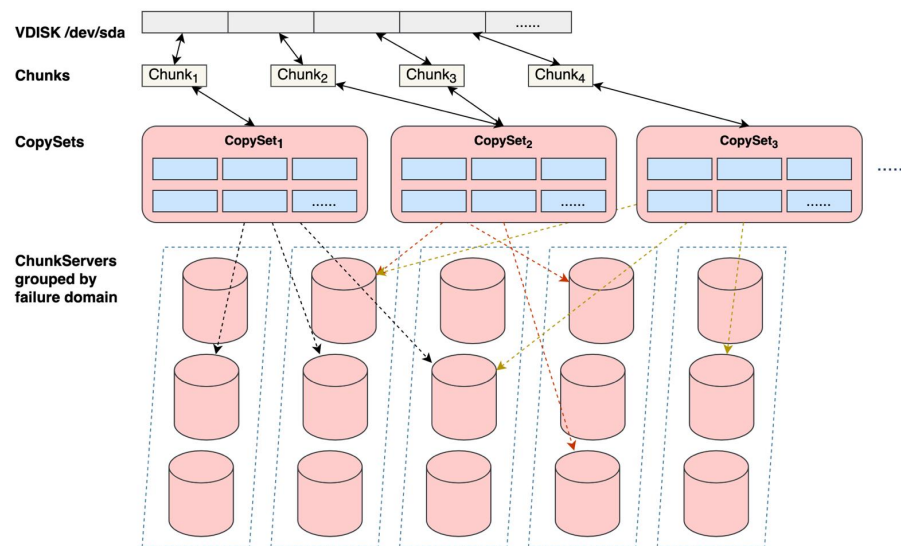
Curve metadata organization

Curve maps virtual block devices to files

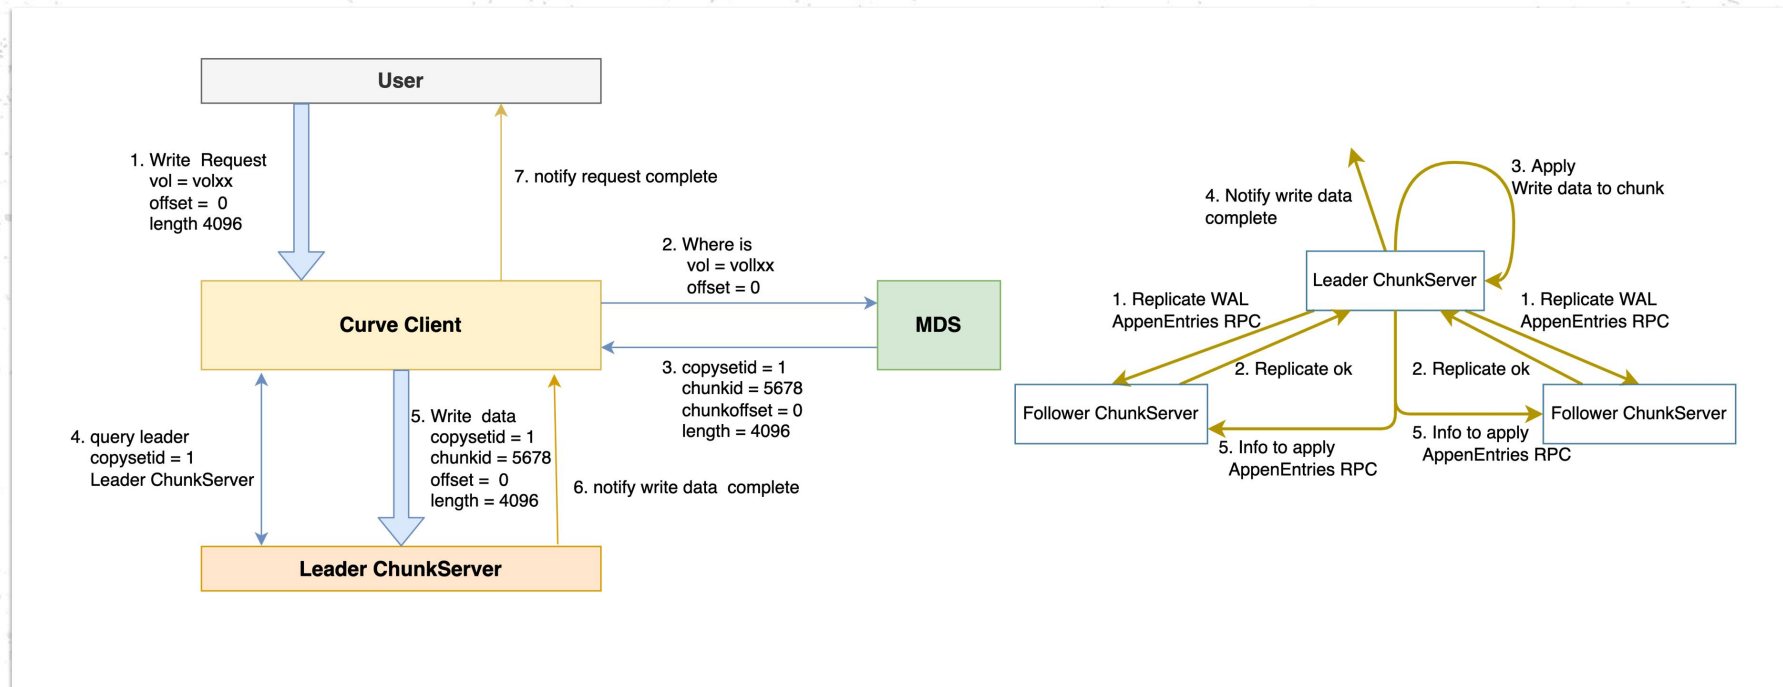
Each file contains chunks scattered across storage nodes in the cluster

Chunkservers are grouped by failure domains

Nodes in a COPYSET belong to different failure domains



CURVE IO data flow



Other performance optimizations

RAFT protocol

Zero data copy

pre-created file pool

yanjiangxu commented on 10 May 2017 · edited

The behavior in original flow:

- 1) In the case of 3 copies, the client sends the write to the primary, and then the primary forwards to the other two replicas. The primary must wait for the commit until all the replicas from completion. The real write IO latency is the longest latency among three copies.
- 2) When the data distribution is not uniform, the osd node which owns hottest data will become the bottleneck of total latency.

The behavior in my enhancement:

commit_majority function is not required to wait for all copies of commit to complete: as long as the majority of the commit is completed, the cluster would notify the client to complete.

As in my implementation:

We also follow the scenario: When the number of copies is less than min_size, ceph does not provide read and write capabilities. In order to ensure security, you must ensure that the number of commit to majority.

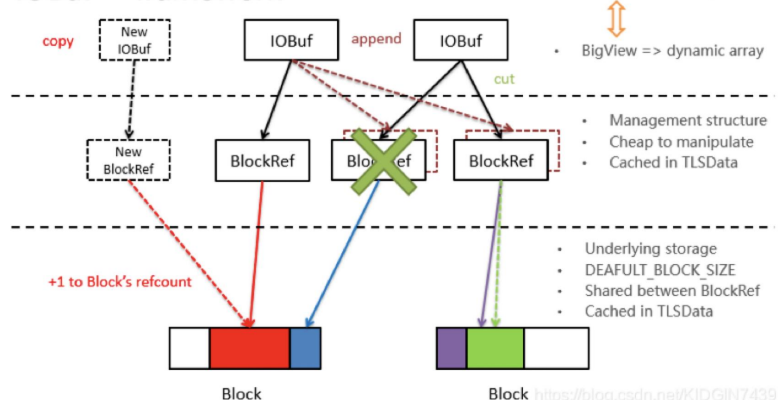
And I defined two important variables:

```
allow_uncommitted_replicas = pool_min_size - 1,
majority = pool_default_size - allow_uncommitted_replicas.
```

Test case:

- 1) 3 copies, pool_default_size = 3, pool_min_size = 2, majority = 2, allow_uncommitted_replicas=1
- 2) The average latency decreased by 11.5%;
- 3) Long tail (99.9%) latency decreased by 71.59%;
- 4) IOPS increased by 13%.

IOBuf -- framework



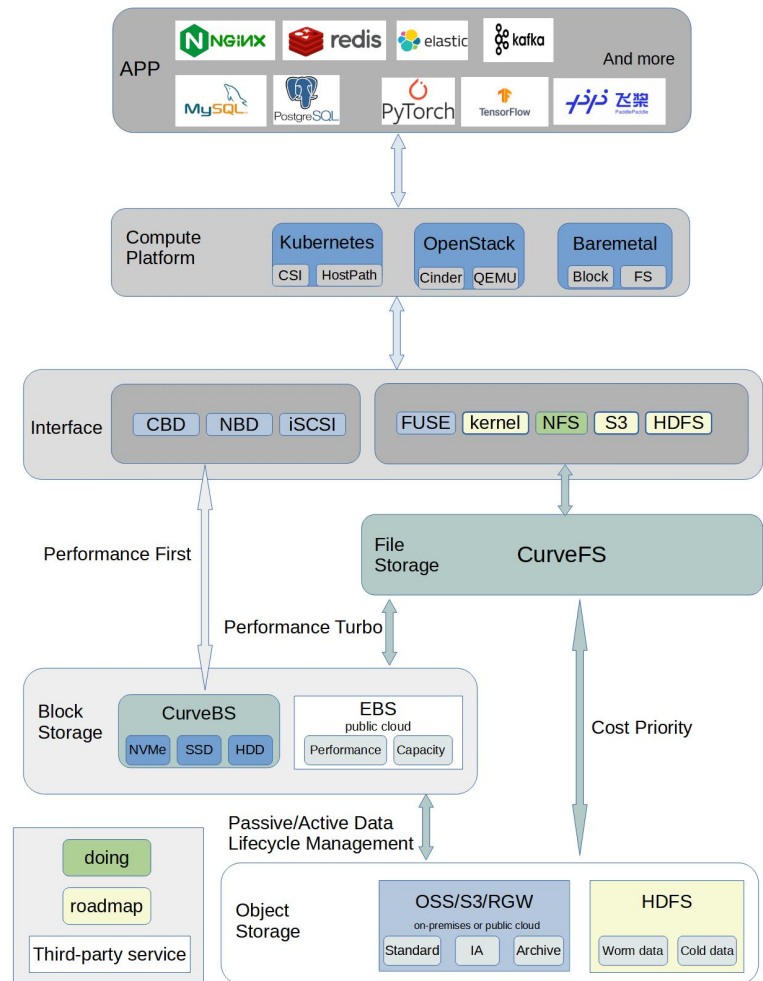
CURVE file system

File service middleware

Upper-layer applications can access file data in storage pools through interfaces such as NFS/HDFS/Posix api

Manage multiple types of storage (object storage, HDFS storage, Elastic block storage)

Support both on-premise and public clouds



Agenda

Why develop storage

Design objectives

Achievements in CURVE

Key designs used by CURVE

✓ CURVE Roadmap

✓ CURVE Community

✓ You can take away



CURVE Roadmap

- Features
 - Supports RDMA and SPDK
 - Further reduce I/O latency and improve throughput
 - Support data lifecycle management
 - Curve block devices can be used as the underlying storage of Curve file systems
- System operation and maintenance
 - Further improve tools for installation/monitoring/maintenance
 - dashboard
- Cloud Native
 - Using CURVE file system in Kubernetes



CURVE community

The improvement and development of Curve cannot be achieved without the support and participation of community contributors.

Project address <https://github.com/opencurve/curve>

Release cycle: a major release every six months, and a minor release every 1-2 months

Follow up on CURVE progress: CURVE meetings are held every 2 weeks to explain CURVE progress and discuss related issues

Submit bugs and Suggestions:
<https://github.com/openCURVE/CURVE/issues>

Participated in the communication and discussion of Curve project: wechat group, qr code on the right hand side



You can take away

Challenges faced

- With the development of hardware and software, more and more stateful applications require to run on the clouds
- High availability/reliability
- Easy to operation and maintenance
- Cloud Native

CURVE solution

- RAFT protocol; Copyset allocation algorithm with topology-based failure domain to provide high availability/reliability
- Zero copy; Data stripe; RDMA to improve performance
- Management and monitor tools
- Support CSI driver for upper cloud native applications access



THANKS